Generators in Python

Square Processing

Let's say you need to do something with square numbers.

```python
def fetch_squares(max_root):
    squares = []
    for n in range(max_root):
        squares.append(n**2)
    return squares


MAX = 5
for square in fetch_squares(MAX):
    do_something_with(square)
```

This works. But...

Maximum MAX

What if MAX is not 5, but 10,000,000? Or 10,000,000,000? Or more?

What if you aren't doing arithmetic to get each element, but making a truly expensive calculation? Or making an API call? Or reading from a database?

Now your program has to wait... to create and populate a huge list... before the second for-loop can even START.

Lazily Looping

The solution is to create an iterator to start with, which lazily computes each value just as it's needed. Then each cycle through the loop happens just in time.

The Iterator Protocol

Here's how you do it in Python:

```python
class Squares:
    def __init__(self, max_root):
        self.max_root = max_root
        self.root = 0
    def __iter__(self):
        return self
    def __next__(self):
        if self.root == self.max_root:
            raise StopIteration
        value = self.root ** 2
        self.root += 1
        return value

for square in Squares(5):
    print(square)
```

There's got to be a better way

Good news. There's a better way.

It's called the **generator**. You're going to love it!

- Sidesteps potential memory bottlenecks, to greatly improve scalability and performance
- Improves real-time responsiveness of the application
- Can be chained together in clear, composable code patterns for better readability and easier code reuse
- Provides unique, valuable mechanisms of encapsulation. Concisely expressive and powerfully effective coding
- A key building block of the async services in Python 3

Yield for Awesomeness

A generator looks just like a regular function, except it uses the `yield` keyword instead of `return`.

```
>>> def gen_squares(max_root):
...     for n in range(max_root):
...         yield n * n
...
>>> for square in gen_squares(5):
...     print(square)
...
0
1
4
9
16
```

Generator Functions & Objects

The function with `yield` is called a **generator function**.

The object it returns is called a **generator object**.

```
>>> def gen_squares(max_root):
...     for n in range(max_root):
...         yield n ** 2
...
>>> squares = gen_squares(5)
>>> type(squares)
<class 'generator'>
```

Pop quiz

Create a new file called `gensquares.py`. Type this in and run it:

```python
def gen_squares(max_root):
    for n in range(max_root):
        yield n ** 2
squares = gen_squares(5)
for square in squares: print(square)
```

It should print:

```
0
1
4
9
16
```

When done, give a thumbs up, comment out the `for` loop, and replace it with `print(next(squares))` repeated several times. What does that do?

The next() thing

```
>>> squares = gen_squares(5)
>>> next(squares)
0
>>> next(squares)
1
>>> next(squares)
4
>>> next(squares)
9
>>> next(squares)
16
>>> next(squares)
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
StopIteration
```

Future-proofing "next"

```python
def gen_up_to(limit):
    n = 0
    while n <= limit:
        yield n
        n += 1


it = gen_up_to(10)

# Works in Python 3 only
it.__next__()
# Works in Python 2 only
it.next()
# Works in Python 2, 3, 4, ...
next(it)
# next() also lets you supply a default value
next(it, None)
```

Multiple Yields

You can have more than one yield statement.

```python
>>> def myitems(limit):
...     for n in range(limit):
...         yield n*n
...     yield "All done"
...
>>> for item in myitems(3):
...     print(item)
...
0
1
4
All done
```

Lab: Generators

Lab file: `generators/generators.py`

- In labs/py3 for 3.x; labs/py2 for 2.7

- When you are done, give a thumbs up...

- ... and then do `generators/generators_extra.py`

NOTE: If the test fails saying it sees `<class 'generator'>`, but expected `<type 'generator'>` - or the other way around - check your Python version.

Iterators

An **iterator** is an object that produces the values in a sequence, one at a time.

```
>>> numbers = [ 0, 1, 2, 3 ]
>>> # iter() calls numbers.__iter__()
... it = iter(numbers)
>>> for n in it: print(n)
0
1
2
3
# second time: no output
... for n in it: print(n)
>>>
```

You can only use it once.

Iterable (vs. "iterator")

An object is said to be **iterable** if it provides a way to spawn new iterators, by passing it to `iter()`.

Normally you never need to do this. `for` loops do it automatically.

```
>>> # A good working definition: an object is iterable
... # if you can pass it to more than one for loop.
... names = [ "Joe", "Tina", "Abdul"]
>>> for name in names: print(name)
Joe
Tina
Abdul
>>> for name in names: print(name)
Joe
Tina
Abdul
```

GenObjs are iterators (not iterable)

```python
def gen_squares(max_root):
    for n in range(max_root):
        yield n * n
```

```python
>>> def gen_squares(max_root):
...     for n in range(max_root):
...         yield n * n
>>> squares = gen_squares(2)
>>> for square in squares: print(square)
0
1
>>> for square in squares: print(square)
... # no output
>>>
```

Iterator Protocol

Any object in Python can be an iterator. It just needs to define proper `__iter__` and `__next__` methods.

```python
class Squares:
    def __init__(self, max_root):
        self.max_root = max_root
        self.current_base = 0
    def __iter__(self):
        return self
    def __next__(self):
        if self.current_base == self.max_root:
            raise StopIteration
        value = self.current_base ** 2
        self.current_base += 1
        return value

for square in Squares(5):
    print(square)
```

We call this the *iterator protocol*.

Dictionary Views & Iteration

Here's a Python 3 dictionary:

```
>>> calories = {
...     "apple": 95,
...     "slice of bacon": 43,
...     "cheddar cheese": 113,
...     "ice cream": 15, # You wish!
... }
>>> items = calories.items()
>>> type(items)
<class 'dict_items'>
>>> hasattr(items, '__next__')
False
>>> hasattr(items, '__iter__')
True
```

What is returned by .items()?

A *dictionary view* object.

Quacks like a dictionary view if it supports three things:

- `len(view)` returns the number of items
- `view` is iterable (e.g., defines `__iter__`)
- `(key, value) in view` returns `True` if that pair is in the dictionary; else, `False`.

Iterable Views

A view is iterable, so you can use it in a for loop:

```
>>> for food, count in calories.items():
...     print("{:.<20s} {:<d} cal".format(food, count))
...
ice cream........... 15 cal
slice of bacon...... 43 cal
apple............... 95 cal
cheddar cheese...... 113 cal
```

Dynamically updates

A view dynamically updates, even if the source dictionary changes:

```
>>> items = calories.items()
>>> len(items)
4
>>> calories['orange'] = 50
>>> len(items)
5
>>> ('orange', 50) in items
True
```

Other methods

There are two other methods on dictionaries, called `.keys()` and `.values()`. They also return views.

```
>>> foods = calories.keys()
>>> counts = calories.values()
>>> 'yogurt' in foods
False
>>> 100 in counts
False
>>> calories['yogurt'] = 100
>>> 'yogurt' in foods
True
>>> 100 in counts
True
```

Benefits

Views improve over regular iterators:

- Providing a more dynamic view of the dictionary's contents

- Let you pass dict contents to caller and know it won't be modified

- Of course, more scalable/performant than a list of (key, value) pairs

What about Python 2?

All the above was for Python 3. Here's how it works in 2:

- `calories.items()` returns a list of `(key, value)` tuples.
    - So if it has 100,000 entries...

- `iteritems()`: returns an iterator over the key-value tuples
- `viewitems()`: which returned a view

`iteritems` is basically obsoleted by `viewitems`, but most people don't realize this yet.

Obsolete methods

In Python 3, what used to be called `viewitems()` was renamed `items()`, and the old `items()` and `iteritems()` went away.

If you still need an actual list in Python 3, you can just say `list(calories.items())`